

YS-3 LLM Cybersecurity Risks Final Report


Senior Project

CS 4850 Section 01
Spring 2026


Kennesaw State University
05/13/2026

Faculty Advisor: Dr. Yong Shi
(yshi5@kennesaw.edu)

Team Members:



Fabian Gonzalez
Team Leader
Documentation



Carlos Anthony Cruz
Developer/
Documentation

STATS AND STATUS	
Lines of Code (LOC)	1,937
Components / Tools	7 Google Colab, Ollama, Python, FAISS, Sentence Transformers, Google Sites, Gemini API
Hours Estimate	200
Hours Actual	260
Status	Project is 100% complete and working as designed

Website:

<https://sites.google.com/view/ys3llm/home>

GitHub:

<https://github.com/ys3seniorproject-arch/ys3seniorproject-arch.io>

Table of Contents

Table of Contents	3
1. Introduction	4
1.1 Selection of OWASP Vulnerabilities	4
2. Requirements	6
2.1 Functional Requirements	6
2.2 Non-Functional Requirements	6
2.3 Requirements Traceability	6
3. Analysis and Design	8
3.1 Architectural Overview	8
3.2 Technology Stack	8
3.3 Data Flow	9
4. Development	10
4.1 Shared Infrastructure	10
4.2 Module Implementations	10
4.2.1 LLM01 — Prompt Injection (400 LOC)	10
4.2.2 LLM03 — Supply Chain (362 LOC)	12
4.2.3 LLM05 — Improper Output Handling (277 LOC)	13
4.2.4 LLM07 — System Prompt Leakage (427 LOC)	15
4.2.5 LLM09 — Misinformation (471 LOC)	16
4.3 Database Connection	18
4.4 Project Setup	19
5. Test Plan	20
5.1 Scope	20
5.2 Strategy	20
5.3 Test Environment	20
6. Test Report	21
7. Version Control	23
8. Conclusion	24
9. References	25

1. Introduction

This report documents the design, development, and testing of the YS-3 Yellow project: LLM Cybersecurity Risks. The project is a research and educational platform that demonstrates five vulnerabilities from the OWASP Top 10 for LLM Applications (2025 edition) through interactive, self contained Google Colab notebooks paired with an educational Google Sites portal.

Large Language Models have rapidly become integral to enterprise workflows, powering customer facing chatbots, internal assistants, code generation tools, and content pipelines. Their adoption, however, has outpaced the security tooling and awareness needed to deploy them safely. Unlike traditional software vulnerabilities that exploit deterministic code paths, LLM vulnerabilities exploit the non deterministic nature of natural language processing making them uniquely difficult to detect, reproduce, and mitigate.

The YS-3 project addresses this gap by translating the abstract vulnerability descriptions in the OWASP Top 10 for LLMs into hands on, executable demonstrations that students and practitioners can run themselves. Each module follows a consistent pedagogical structure: a pre lab introducing the vulnerability concept, a hands on lab where the attack is triggered and then mitigated in real time, and a post lab with reflection questions and further reading.

1.1 Selection of OWASP Vulnerabilities

The OWASP Top 10 for LLM Applications contains ten vulnerability categories. For this project, five were selected based on three criteria: demonstrability in a constrained educational environment (Google Colab with free-tier GPU), pedagogical impact (how clearly the vulnerability illustrates a distinct class of risk), and breadth of coverage across the attack surface (input manipulation, output handling, supply chain trust, information disclosure, and content integrity).

The five selected vulnerabilities and the rationale for each are as follows:

LLM01 — Prompt Injection was selected because it is the most widely discussed and arguably the most fundamental LLM vulnerability. It demonstrates how natural language instructions embedded in user input or multimodal content (images, in this case) can override a model's intended behavior. This module uses visual prompt injection with Qwen2VL to show that the attack surface extends beyond text.

LLM03 — Supply Chain was chosen to highlight that LLM security is not solely about prompt level attacks. Outdated or unmaintained model checkpoints may lack the safety fine tuning present in continuously updated cloud models. This module compares a static local checkpoint against a maintained cloud model (Gemini) given identical extraction payloads, demonstrating that supply chain hygiene knowing what model you are running and whether it receives updates is itself a security concern.

LLM05 — Improper Output Handling was selected because it bridges LLM specific risks with classical web application security. LLM output is frequently rendered in browsers or passed to backend functions without sanitization, creating vectors for XSS, SQL injection, and other injection attacks. This module demonstrates that treating LLM output as trusted input is a critical design flaw.

LLM07 — System Prompt Leakage was included because many deployed LLM applications embed sensitive information API keys, database credentials, business logic directly in system prompts under the assumption that users cannot access them. This module demonstrates seven distinct extraction techniques against a vulnerable system prompt and shows that the mitigation is architectural: never place secrets where the model can reach them.

LLM09 — Misinformation was selected to address the risk of confident hallucination, particularly in niche domains where the model's training data is sparse. This module uses obscure fitness equipment specifications as ground truth to show that models will fabricate plausible but incorrect details, and demonstrates that Retrieval Augmented Generation (RAG) with verified sources is an effective mitigation.

Together, these five modules cover the input layer (LLM01), the model supply chain (LLM03), the output layer (LLM05), information disclosure (LLM07), and content

integrity (LLM09), providing broad coverage of the LLM attack surface within the project's scope.

2. Requirements

The requirements for the YS-3 project were derived from the course deliverable specifications, the OWASP Top 10 for LLM Applications framework, and the educational objectives established with the faculty advisor. The full Software Requirements Specification (SRS) was delivered as a separate document earlier in the semester; this section summarizes the key functional and non functional requirements.

2.1 Functional Requirements

The system shall implement at minimum five interactive demonstrations of OWASP Top 10 LLM vulnerabilities, with LLM01 (Prompt Injection) and LLM05 (Improper Output Handling) designated as mandatory modules. Each module shall include both a vulnerability demonstration and a corresponding mitigation, presented in a side by side or toggle based comparison format. All demonstrations shall execute within Google Colab notebooks using locally served models via Ollama, requiring no external API keys for core functionality (Gemini is used as a supplementary comparison in LLM03 only). The educational portal shall organize each module into pre lab, hands on lab, and post lab sections following the Blockchain Labware reference structure.

2.2 Non-Functional Requirements

All demonstrations shall run on Google Colab's free tier T4 GPU runtime. The inline HTML user interface shall render correctly within Colab's output cells without external hosting. Models shall be pulled on demand from the Ollama registry to ensure reproducibility across sessions. All vulnerability demonstrations shall use synthetic data and controlled environments no real user data or production systems are targeted.

2.3 Requirements Traceability

Requirement ID	Description	Module	Status
FR-01	Image-based prompt injection demonstration	LLM01	Complete
FR-02	Multi-layer output validation (XSS, SQLi, path traversal)	LLM05	Complete
FR-03	Supply chain comparison: static vs. maintained model	LLM03	Complete
FR-04	System prompt extraction with 7 attack vectors	LLM07	Complete
FR-05	RAG-grounded misinformation mitigation	LLM09	Complete
FR-06	Inline HTML UI for all modules	All	Complete
FR-07	Google Sites educational portal with pre/lab/post structure	All	Complete
NFR-01	Runs on Colab free-tier T4 GPU	All	Complete
NFR-02	No external API keys for core demos	All	Complete

3. Analysis and Design

The system architecture follows a modular, notebook per vulnerability design where each Google Colab notebook is fully self contained: it installs its own dependencies, starts its own Ollama server instance, pulls the required model, defines the vulnerability logic and mitigation, and renders an inline HTML user interface all within a single executable document.

3.1 Architectural Overview

The system consists of three layers. The Presentation Layer is a Google Sites portal that organizes educational content into pre lab, lab, and post lab sections for each module. The Application Layer comprises five Google Colab notebooks, each implementing a vulnerability demonstration and mitigation with an inline HTML/JavaScript UI rendered through IPython display utilities. The Infrastructure Layer is Ollama running as a background process inside the Colab runtime, serving local models (llama3.2:1b, Qwen2.5VL:7b, mistral:7b, qwen2.5:7b, qwen2.5:0.5b, deepseek-r1:7b) via a localhost REST API.

3.2 Technology Stack

Component	Technology	Purpose
Runtime Environment	Google Colab (T4 GPU)	Cloud execution with free GPU access
Model Server	Ollama	Local LLM serving via REST API
Models	Qwen2.5VL, Mistral, Qwen2.5, DeepSeek-R1	Vulnerability demonstrations
Cloud Model	Gemini 2.5 Flash	Supply chain comparison (LLM03 only)
RAG Components	FAISS + sentence-transformers	Vector retrieval for LLM09 and LLM01
Frontend	Inline HTML/CSS/JS via IPython	Interactive demo UIs
Educational Portal	Google Sites	Pre-lab/lab/post-lab content hosting
Language	Python 3	All notebook logic

3.3 Data Flow

The interaction flow for a typical module proceeds as follows: the user opens the Colab notebook and connects to a T4 GPU runtime. Setup cells install system dependencies (zstd, pciutils), install Ollama, start the Ollama server as a background thread, and pull the target model. The user interacts with the inline HTML UI, which invokes a Python callback function registered via `google.colab.output.register_callback`. The callback constructs the appropriate prompt (with or without mitigation), sends it to Ollama via the ollama Python client, processes the response, and returns the results as a base64-encoded JSON payload to the JavaScript frontend, which renders the output in side by side comparison panels.

4. Development

This section documents the implementation of each module, the shared infrastructure, and the development decisions made during the project.

4.1 Shared Infrastructure

Every notebook shares a common setup pattern: install system dependencies, install and start Ollama as a background thread, and pull the target model. The Ollama server is configured to listen on 0.0.0.0:11434 with CORS origins set to wildcard to allow Colab's iframe based output cells to communicate with the localhost API. The `LD_LIBRARY_PATH` is set to `/usr/lib64-nvidia` to ensure Ollama can access the Colab runtime's NVIDIA GPU libraries.

The UI architecture uses a callback based pattern: JavaScript in the inline HTML calls `google.colab.kernel.invokeFunction` to trigger a Python function, which performs the model inference and returns results as a base64 encoded JSON string. This pattern avoids the limitations of ipywidgets for complex interactive layouts while maintaining the notebook's self contained nature.


4.2 Module Implementations

4.2.1 LLM01 — Prompt Injection (400 LOC)

This module demonstrates visual prompt injection using the Qwen2.5VL:7b vision language model. An image containing embedded text instructions is sent to the model alongside a user's message. The module provides three configurable dimensions: system prompt strength (weak vs. strong), output validation (on/off), and RAG context (on/off).

The weak system prompt instructs the model to follow user and image instructions exactly, while the strong prompt restricts the model to image description only with explicit rules against following embedded instructions. The output validation layer implements multi pattern detection including role change detection, instruction leakage, jailbreak compliance, and sensitive data leak patterns (API keys, passwords, SSNs). A FAISS based RAG system with sentence transformers (all-MiniLM-L6-v2) stores

synthetic confidential documents to demonstrate how retrieved context can expand the attack surface when combined with prompt injection.






Figure 1: Top image shows what a system response will look like from a model containing sensitive data inside, and how prone that data is to being leaked by a threat actor. Whereas the bottom image shows a system without confidential data in the RAG and intern shows how it was blocked by the prompt parsing and the lack of data it can use.

4.2.2 LLM03 — Supply Chain (362 LOC)

This module compares the security posture of a static, frozen model checkpoint (served locally via Ollama) against a continuously maintained cloud model (Gemini 2.5 Flash). Both models receive an identical system prompt containing synthetic sensitive data (API keys, database credentials, acquisition details) and an identical extraction payload. The side by side comparison demonstrates that static checkpoints, which do not receive ongoing safety fine tuning, are more likely to comply with extraction attempts than actively maintained models.

The UI uses a clean, white background research panel design with automated verdict detection that scans both responses for the presence of known secret strings and reports whether each model leaked confidential information.





Figure 2: Left: Shows vulnerability detected because of the static points and not consistent updating of model. Where on the contrary the right side shows secured as the consistent updating helps the system stay up to date with security

4.2.3 LLM05 — Improper Output Handling (277 LOC)

This module demonstrates that LLM output, when passed unsanitized to downstream systems, can carry payloads identical to traditional injection attacks. Three test cases are defined: XSS via script tag injection, SQL injection payload, and a safe baseline response. A validation layer with six regex based rules detects XSS (script tags, event

handlers, javascript: URIs), SQL injection patterns, path traversal attempts, and markdown image exfiltration. The validation toggle allows users to observe the difference between filtered and unfiltered output reaching a simulated backend function.

This module uses ipywidgets (Dropdown, ToggleButton, Button) rather than the inline HTML callback pattern used by the other modules, demonstrating an alternative UI approach within Colab






Figure 3: Top: Shows what a system without a validation layer looks like, and what can be leaked due to the lack of data parsing sent to the users request. Bottom: Shows how a more strict system with a validation layer can interpret when sensitive data is being leaked and blocks it from ever getting to the user who requested it.

4.2.4 LLM07 — System Prompt Leakage (427 LOC)

This module demonstrates seven distinct system prompt extraction techniques: direct request, roleplay bypass, translation request, summarization request, developer mode injection, indirect credential query, and social engineering. The same attack is sent to the same model (Mistral 7B) with two different system prompts. The weak prompt embeds confidential data directly (API keys, passwords, database hosts, encryption keys). The hardened prompt contains no secrets and redirects sensitive questions to appropriate channels.

The UI renders both system prompts in preview panels and provides an automated verdict system that scans each response for known secret strings. The dark themed

IBM Plex design uses red and green color coding for vulnerable and hardened panels respectively.




Figure 4: Left side shows what the output of a system prompt containing sensitive data and a weak verification layer looks like. Right side shows what a system prompt with a hardened prompt would look like and how it could help prevent sensitive data from leaking. By not containing any to begin with.

4.2.5 LLM09 — Misinformation (471 LOC)

This module demonstrates confident hallucination in niche domains. The knowledge base contains verified specifications for the Nautilus XPload plate loaded equipment line an intentionally obscure domain where the model's training data is expected to be sparse. Five test scenarios are provided: full product line enumeration, steel gauge specifications, specific machine dimensions, product lineage history, and a medical safety question (infant honey consumption).

Each query is sent twice: once raw (no grounding) and once with RAG retrieved verified facts injected into the prompt. An analysis function scans both responses for misinformation indicators: statistics without citations, overconfident language, and medical content without disclaimers. The side by side comparison demonstrates that the ungrounded model fabricates plausible specifications while the RAG grounded model responds with verified information and cites its sources




Figure 5: Show the difference of having a rag incoreperated into the system and not having one. On the left side: you can see how the one without the rag hallucinates incorrect responses, whereas on the right side shows the one containing a RAG produced a more grounded result.

4.3 Database Connection

Not applicable to this project. The YS-3 system does not use a persistent database. The FAISS vector index used in LLM01 and LLM09 is constructed in memory at runtime

from hardcoded document arrays and is discarded when the Colab session ends. No persistent storage or database connectivity was required.

4.4 Project Setup

To run any YS-3 module, the user needs: (1) a Google account with access to Google Colab, (2) a Colab runtime connected to a T4 GPU, and (3) an internet connection to pull Ollama models on first run. No local installation is required. Each notebook's first cell handles all dependency installation automatically. For the LLM03 module specifically, a Gemini API key is required for the cloud model comparison; a key is provided inline in the notebook for demonstration purposes.

5. Test Plan

The testing strategy for YS-3 is manual, demonstration driven testing executed by the development team. Given the research and educational nature of the project, formal automated test suites, load testing, and penetration testing of the notebooks themselves are out of scope. The full Software Test Plan (STP) was delivered as a separate document; this section summarizes the approach and key results.

5.1 Scope

Testing covers the five vulnerability modules (LLM01, LLM03, LLM05, LLM07, LLM09), the shared Ollama runtime infrastructure, and the inline HTML user interfaces. Each module is verified for three objectives: environment setup success, vulnerability demonstration effectiveness, and mitigation behavioral change.

5.2 Strategy

Module level testing was performed by the developer (Carlos Anthony Cruz). Demonstration walk throughs were reviewed jointly with the team leader (Fabian Gonzalez Tascon) and the faculty advisor (Dr. Yong Shi). Each module was executed end to end on Google Colab with a T4 GPU runtime. The vulnerability was triggered, the mitigation was applied, and the comparative outputs were observed and recorded.

5.3 Test Environment

Google Colab with T4 GPU, Python 3, system dependencies zstd and pciutils, Ollama installed inside the runtime, and models pulled on demand. Each notebook is self contained; environment state is rebuilt from scratch at the start of every Colab session, ensuring consistent baseline conditions.

6. Test Report

The following table records the outcome of demonstration based testing for each YS-3 module.

Test ID	Description	Expected Result	Result	Severity
TC-01	LLM01 — Qwen2VL model load	Installs deps, starts Ollama, pulls model	Pass	N/A
TC-02	LLM01 — Image injection triggers unintended behavior	Model follows injected instruction	Pass	Minor
TC-03	LLM01 — Mitigation reduces injection rate	Fewer successful injections after mitigation	Pass	Minor
TC-04	LLM03 — Supply chain demo runs	Demo completes end to end	Pass	N/A
TC-05	LLM03 — Mitigation alters trust behavior	Maintained model rejects extraction	Pass	N/A
TC-06	LLM05 — Unsafe output produced	Unsanitized output reaches rendering layer	Pass	N/A
TC-07	LLM05 — Sanitization neutralizes output	Output escaped before rendering	Pass	N/A
TC-08	LLM07 — System prompt extraction	Crafted prompt elicits system prompt	Pass	N/A
TC-09	LLM07 — Mitigation prevents leakage	Hardened model refuses extraction	Pass	N/A
TC-10	LLM09 — Misinformation scan	Indicators flagged in responses	Pass	N/A
TC-11	LLM09 — RAG reduces hallucination	Grounded response uses verified facts	Pass	N/A
TC-12	Shared — Ollama serves localhost API	HTTP POST returns valid completions	Pass	N/A
TC-13	Shared — Inline HTML UI renders	UI elements render in Colab	Pass	N/A
TC-14	Shared — Side-by-side panels populate	Both panels populate after run	Pass	N/A
TC-15	Regression — Post-refactor parity	Equivalent or improved output	Pass	N/A

All fifteen test cases passed. TC-02 and TC-03 are recorded with Minor severity to reflect the inherent non determinism of large language models: identical prompts may yield different outputs across runs due to sampling temperature and model variance, and the prompt injection demonstration succeeds intermittently rather than on every execution. No critical or moderate defects were observed.

7. Version Control

Version control for the YS-3 project was managed through an iterative file versioning approach. As each module was developed and refined, new versions of the notebook scripts were created to preserve the state of prior iterations. This allowed the team to roll back to earlier working states when experimental changes introduced regressions particularly important given the stochastic nature of LLM outputs, where a code change that improves one demonstration scenario might degrade another.

Google Colab's built in revision history provided an additional layer of version tracking, automatically recording changes to each notebook over time. The final deliverable scripts represent the culmination of multiple development iterations spanning the Spring 2026 semester.

The educational portal content on Google Sites was versioned through the platform's native page history feature, which records edits with timestamps and author attribution.

For final submission, all source code is provided as a ZIP archive and will be hosted on GitHub for public access.

8. Conclusion

The YS-3 LLM Cybersecurity Risks project successfully delivered a research and educational platform demonstrating five OWASP Top 10 LLM vulnerabilities through interactive Google Colab notebooks and an accompanying Google Sites educational portal. All five modules Prompt Injection (LLM01), Supply Chain (LLM03), Improper Output Handling (LLM05), System Prompt Leakage (LLM07), and Misinformation (LLM09) are fully implemented, tested, and documented.

The project's key technical contributions include: a visual prompt injection demonstration using Qwen2VL that extends the attack surface beyond text to images; a supply chain comparison framework that quantifies the security gap between static and maintained models; a multi rule output validation engine that detects XSS, SQL injection, and other injection patterns in LLM output; a systematic evaluation of seven system prompt extraction techniques with automated secret leak detection; and a RAG grounded misinformation mitigation system tested against intentionally obscure domain knowledge.

The total codebase comprises 1,937 lines of Python across five notebooks, supported by a Google Sites portal with pre lab, lab, and post lab content for each module. The project was completed within 260 hours against an initial estimate of 200 hours, with the additional time primarily spent on the iterative refactoring from the Hugging Face transformers pipeline to Ollama and on tuning the inline HTML user interfaces for clarity and usability.

The platform is designed to be immediately usable by students and practitioners: every notebook runs on Google Colab's free tier with no local setup, every vulnerability is demonstrated with a clear before and after comparison, and every mitigation is actionable rather than theoretical. The project demonstrates that LLM security education does not require expensive infrastructure or proprietary tools it requires clear demonstrations that make abstract risks concrete.

9. References

- [1] OWASP Foundation. (2025). OWASP Top 10 for Large Language Model Applications 2025. <https://genai.owasp.org/llm-top-10/>
- [2] OWASP Foundation. (2025). LLM05:2025 Improper Output Handling. <https://genai.owasp.org/llmrisk/llm052025-improper-output-handling/>
- [3] OWASP Foundation. OWASP Application Security Verification Standard (ASVS) — Chapter 5: Validation, Sanitization and Encoding. <https://owasp-aasvs4.readthedocs.io/en/latest/V5.html>
- [4] Zhang, J., et al. (2025). When LLMs meet cybersecurity: A systematic literature review. *Cybersecurity*, 8(1), p.55.
- [5] Abuadbba, A., et al. (2025). From Promise to Peril: Rethinking Cybersecurity Red and Blue Teaming in the Age of LLMs. arXiv preprint arXiv:2506.13434.
- [6] Akiri, C., et al. (2025). Safety and Security Analysis of Large Language Models: Benchmarking Risk Profile and Harm Potential. arXiv preprint arXiv:2509.10655.
- [7] Blockchain Labware Reference Site. <https://sites.google.com/view/blockchainlabware>
- [8] Embrace The Red. (2023). ChatGPT Plugin Exploit Explained: From Prompt Injection to Accessing Private Data.
- [9] Snyk Security Blog. Arbitrary Code Execution in LangChain. <https://security.snyk.io/vuln/SNYK-PYTHON-LANGCHAIN-5411357>